

Homework 4: Unix Signals and (more) File I/O

Due: Monday 28 October 2013, 23:59:00 (11:59 PM) Pacific USA time zone.

Points on this assignment: 205 points with 20 bonus points available.

Work submitted late will be penalized as described in the course syllabus. You must submit your work twice for this and all other homework assignments in this class. Ecampus wants to archive your work through Blackboard and EECS needs you to submit through TEACH to be graded. If you do not submit your assignment through TEACH, it cannot be graded (and you will be disappointed with your grade). Make sure you submit your work through [TEACH](#). Submit your work for this assignment as a *single tar.bzip file* through TEACH. The same single tar.bzip file should also be submitted through Blackboard.

Place all of the files you produce for this assignment in a single directory, called `Homework4`.

In this assignment, you will be working with Unix signals and file I/O system calls. This assignment is a significant step up in complexity from previous homework assignments. This assignment will also be only in C, with a good bit of Makefile automation mixed in. If you look through some of the C code on the previous homeworks, you can find some C code fragments that may be helpful in this assignment. You have 2 weeks to complete this assignment. However, I urge you to not delay beginning it. Starting this soon after or concurrently with homework #3 should work very well.

In **all** your source files (including the Makefile), you need to have 4 things at the top of every file as comments:

1. Your name
2. Your email address (ONID or engineering)
3. The class name and section (this is CS311-400)
4. The assignment number (this is homework #4)

Remember that the programming work in this class is intended to be individual work, not group work.

1. **5 points.** When you are ready to submit your files for this assignment, make sure you submit a *single bzip file*. Review homework #1, problem #1 if you need a refresher on how to do this. If your file is not a single bzip file, you cannot receive points on this assignment.
2. **25 points total.** Write a C program on `eos-class` (or CS311 VM) to send and receive Unix signals. Don't go overboard on this portion of the assignment. You

should be able to easily complete this portion with less than 50 lines of C code. It may take you less time to write the C code than it took me to write the problem description.

You will write an application in C (`sig_demo.c`) that installs 3 different functions as signal handlers. Your C code will handle the following signals:

- `SIGUSR1`
- `SIGUSR2`
- `SIGINT`

It may be tempting to use a single signal handler for all 3 signals, but don't. Create a separate single handler function for each signal.

When your application receives the `SIGUSR1` signal, it should print **(5 points)**:

```
SIGUSR1 has been caught
```

Your application should not exit when the `SIGUSR1` signal is handled.

When your application receives the `SIGUSR2` signal, it should print **(5 points)**:

```
SIGUSR2 has been caught
```

Your application should not exit when the `SIGUSR2` signal is handled.

When your application receives the `SIGINT` signal, it should print **(10 points)**:

```
SIGINT has been caught, terminating the program
```

When your code receives the `SIGINT` signal, your application should exit (after printing the message of course).

You may call `printf()` for these messages within your signal handler. I know that it is not strictly *safe* to call `printf()` within a signal handler, but we will, just this one time.

Once your application has installed the 3 signal handlers, it should send the signals to itself, in this order: `SIGUSR1`, `SIGUSR2`, `SIGINT`. The `getpid()` system call can be your friend here. I want you to use `kill()` not `raise()` to send the signals.

When I run your compiled program, I should see the following

```
$ ./sig_demo
SIGUSR1 has been caught
SIGUSR2 has been caught
SIGINT has been caught, terminating the program
$
```

If you find yourself struggling with this portion of this assignment, you need to review chapter 20 in TLPI and spend some time looking at sections 20.4 and 20.5 and listing 20-1. Again, think in terms of 50 lines of C code.

If you make use of some resources (such as web sites), make sure you reference them in your code.

Put a target in your Makfile to build this code (**5 points**).

3. **140 points total.** This portion of the homework is about reading and writing files. You will need to `stat()` or `fstat()` files, check file permissions, check file time stamps, and perform a `seek()` into a file. You will need more than 50 lines of c code to complete this portion of the assignment.

Write a C program on `eos-class` called `myar`. This program will illustrate the use of file I/O on Unix by maintaining a Unix archive library, in the standard archive format.

Once compiled your program should run in a manner similar to the standard Unix command `ar`. You will need to spend some time looking at the man page for `ar`. However, for this assignment, the following is the syntax your program must support:

```
myar key archive-file [member [...]]
```

where `archive-file` is the name of the archive file to be used, and `key` is one of the following options:

-q	“Quickly” append named files (members) to archive. 20 points . Check the meaning of append in the notes below.
-x	Extract named members. 20 points . Just as the regular <code>ar</code> command, if no member is named on the command line when extracting files, all files are extracted from the archive. The permissions on the extracted files should match permissions on the files before archiving (as described in the notes below).
-t	Print a concise table of contents of the archive. 5 points . The concise table of contents for your application (<code>myar</code>) should match <u>exactly</u> the output from the “ <code>ar t</code> ” command on the same archive file.
-v	Print a verbose table of contents of the archive. 10 points . The verbose table of contents for your application (<code>myar</code>) should match <u>exactly</u> the output from the “ <code>ar tv</code> ” command on the same archive file. See the

	man page on <code>ar</code> .
<code>-d</code>	Delete named files from archive. 50 points. Make sure you read the note below about the <code>-d</code> option and creation of a new file.
<code>-A</code>	Quickly append all “regular” files in the current directory (Except the archive itself). 30 points (of the 140). There is not an option for the Unix <code>ar</code> command that does this.
<code>-w</code>	Extra credit 10 points: for a given timeout (in seconds), add all <u>modified</u> files to the archive (Except the archive itself). See note h. There is not an option for the Unix <code>ar</code> command that does this.

The archive file maintained must use *exactly* the standard format defined in `/usr/include/ar.h`, and in fact may be tested with archives created with the `ar` command. The archive files created the regular `ar` command and those created with your `myar` should be interoperable. Do not copy or in any way modify the `ar.h` include file.

The options listed above are compatible with the options having the same name in the `ar` command, except for the following exceptions: the `-v` and `-t` command take no further argument, and list all files in the archive. `-v` option is short for `-t -v` (or `tv`) on the regular `ar` command. The `-A` and `-w` commands are new options not in the usual `ar` command.

Notes (lots of them and they are all important):

- For the `-q` command `myar` should create an archive file if it doesn't exist, using permissions “666”. For the other commands `myar` reports an error if the archive does not exist, or is in the wrong format. A bad format causes an error statement and your program to exit.
- You will have to use the system calls `stat()` and `utime()` to properly deal with extracting and restoring the proper timestamps. Since the archive format only allows a single timestamp, store the `mtime` and use it to restore both the `atime` and `mtime`. Permissions should also be restored to the original value, subject to `umask` limitations.
- The `-q` and `-A` commands do not check to see if a file by the chosen name already exists. They simply append the file(s) to the end of the archive.
- The `-x` and `-d` commands operate on the first file matched in the archive, without checking for further matches. It is possible for a file name to exist more than once in an archive; use the first one that matches.
- In the case of the `-d` option, you will have to build a new archive file to recover the space. Do this by unlinking the original file after it is opened (or after you’ve completed reading it), and creating a new archive with the original name.
- Since file I/O is slow, do not make more than one pass through the archive file; an issue especially relevant to the multiple member delete case.
- You are required to handle multiple file names as members.

- h) For the `-w` flag, the command will take as long as specified by the timeout argument. You should print out a status message upon adding a new file. This may result in many different copies of the same file in the archive.
- i) Make sure you lookup what a “regular” file is in Unix.
- j) You must have a `Makefile` for this portion of the homework. Your `Makefile` must contain (at least) the following 2 targets:
 1. `all` -- the `all` target should compile and link the final program. The `all` target should also be the default target (which only means that it is the first target in the `Makefile`).
 2. `clean` -- the `clean` target should remove all binary applications (`myar`) and object files (`.o`'s). In addition, you should remove any editor droppings from `emacs`, `vi`, or whatever. The sample `Makefile` from homework #1 included a `clean` target. **Make sure you use this before you bundle all your files together for submission.**

The `Makefile` may (should) contain other targets to build any out of date modules. Again, the sample `Makefile` from Homework1 gives some guidance on this. Targets like `all` and `clean` are sometimes called phony targets because they don't actually produce a file.

When I build your program, I should be able to just type `make` to have it completely build. There are additional requirements for your `Makefile` below.

- k) For **10 points extra credit**, any time a file is added that already exists, remove the old copy from the archive, but only if it is not the same. If identical, do not add the new file. Make sure you clearly comment your meaning of “identical.”
- l) It is not necessary that you use `getopt()` to process `argv` from the command line. If you have a simpler method to use, that is fine. You'll probably get more practice using `getopt()` later in the class.
- m) I have created some sample files that you can use for testing your application. They are all plain text, so you can actually just `cat` the archive file after you've create it to see how it looks and compare it to one created with `ar`. The sample files are: `1-s.txt`, `2-s.txt`, `3-s.txt`, `4-s.txt`, and `5-s.txt`. One of the things you'll note about the test files that some of them have an even number of bytes and others have an odd number of bytes. This will be important to you. The sample files can be found on `eos-class` in:

```
/usr/local/classes/eecs/summer2013/cs311-400/src/Homework4
```

- n) Because I know this is tricky and can cause undue grief, ***you need to carefully read the following 2 web pages and search for the word “even”:*** http://en.wikipedia.org/wiki/Ar_%28Unix%29 and <http://www.unix.com/man-page/opensolaris/3head/ar.h/> One of the other things you learn in this assignment is how to carefully read specifications and man pages. Just take my word for this and carefully read the 2 web pages.
- o) The length of file names used will not exceed 15 characters for this assignment.

- p) Make sure your code compiles before you submit it. Simply being able to compile your code is worth **5 points**.
- q) Because you are *potentially* working with binary files, you should use the `read()` and `write()` system calls, not `fscanf()` and `fprintf()`, for the file I/O. You can use `printf()` for terminal output.
- r) Do yourself a favor and do not try adding binary (non-text) files into your archive file for testing against the regular `ar` command. You'll get some odd messages about a table of contents that won't be very helpful. Test things against `ar` using plain text files. I'll only be testing your code using plain text files.

Since there are quite a few parts to completing this problem in the homework, I suggest that you start small. Start with just creating code that performs the `-t` option on an archive file you've created with the regular `ar`. From there, move to the `-v` option on an archive file. From there, I'd go `-x`, `-q`, `-d`, `-A`, and `-w`, but go in the order that makes a clear progression for you. Remember, the archive files you create with `myar` need to be interoperable with those created with `ar`, so you can easily test one from the other. Having completed homework #3, you already have a head start on completion of this `-t` and `-v` options of the program.

Printing in octal is probably sometime you've not done for quite a while. C does make it pretty easy for you. This web page gets low marks for beauty, but high marks for summarizing all those options for `printf()`:

<http://wpollock.com/CPlus/PrintfRef.htm>. This is a nice page from "The C Book", listed as a recommended reference for C programming:
http://publications.gbdirect.co.uk/c_book/chapter9/formatted_io.html.

If you find yourself struggling with this portion of the assignment, you should find some code fragments from previous assignments (1 and 3 especially) that will give you a good boost. Look for them. If you find yourself puzzling over even and odd file sizes, make sure you read the entire assignment description (including all the many notes).

If you find that your C programming or Makefile skills are still a bit rusty, you may want to peruse some of the tutorials shown at the bottom of the week 3 web page on Blackboard. The book "The C Book" is pretty good and free:

http://publications.gbdirect.co.uk/c_book/

- 4. **35 points total.** Put tests as targets in your Makefile. Since you need to have a number of tests for your code, put them into the Makefile so it is easy for you to run then easily and consistently. This is a simple form of regression testing.
 - a. Put a target called `testq12345` in your Makefile. The `testq12345` target will do the following:
 - i. Remove any file named `ar12345.ar`. If the file `ar12345.ar` does not exist, don't show an error.
 - ii. Remove any file named `myar12345.ar`. If the file `myar12345.ar` does not exist, don't show an error.

- iii. Create a file named `ar12345.ar` using this call:
`ar q ar12345.ar 1-s.txt 2-s.txt 3-s.txt 4-s.txt 5-s.txt`
- iv. Create a file named `myar12345.ar` using this call:
`myar -q myar12345.ar 1-s.txt 2-s.txt 3-s.txt 4-s.txt 5-s.txt`
- v. Compare the files created by `ar` and `myar`.
`diff ar12345.ar myar12345.ar`
- vi. The result of the `diff` command should show no differences.
- b. Put targets in your Makefile called `testq135` and `testq24` which are like the target `testq12345`, but only use the noted subset of files (1, 3, 5, and 2, 4).
- c. Put a target in your Makefile called `testq` that will run the `testq12345`, `testq135`, and `testq24` targets. **5 Points.**
- d. Put a target called `testt12345` in your Makefile. The `testt12345` target will do the following:
 - i. Remove any file named `ar12345.ar`. If the file `ar12345.ar` does not exist, don't show an error.
 - ii. Create a file named `ar12345.ar` using this call:
`ar q ar12345.ar 1-s.txt 2-s.txt 3-s.txt 4-s.txt 5-s.txt`
 - iii. Run the following commands:
`ar t ar12345.ar > ar-ctoc.txt`
`myar -t ar12345.ar > myar-ctoc.txt`
 - iv. Compare the table of contents files by using this command:
`diff ar-ctoc.txt myar-ctoc.txt`
 - v. The result of the `diff` command should show no differences.
- e. Put targets in your Makefile called `testt135` and `testt24` which are like the target `testt12345`, but only use the noted subset of files (1, 3, 5, and 2, 4).
- f. Put a target in your Makefile called `testt` that will run the `testt12345`, `testt135`, and `testt24` targets. **5 Points.**
- g. Put a target called `testv12345` in your Makefile. The `testv12345` target will do the following:
 - i. Remove any file named `ar12345.ar`. The file `ar12345.ar` does not exist, don't show an error.
 - ii. Create a file named `ar12345.ar` using this call:
`ar q ar12345.ar 1-s.txt 2-s.txt 3-s.txt 4-s.txt 5-s.txt`
 - iii. Run the following commands:
`ar tv ar12345.ar > ar-vtoc.txt`
`myar -v ar12345.ar > myar-vtoc.txt`
 - iv. Compare the table of contents files by using this command:
`diff ar-vtoc.txt myar-vtoc.txt`
 - v. The result of the `diff` command should show no differences.
- h. Put targets in your Makefile called `testv135` and `testv24` which are like the target `testv12345`, but only use the noted subset of files (1, 3, 5, and 2, 4).
- i. Put a target in your Makefile called `testv` that will run the `testv12345`, `testv135`, and `testv24` targets. **5 Points.**

- j. Put a target in your `Makefile` called `tests` that will run the `testq`, `testt`, and `testv` targets. With this, you should only need to run “`make tests`” to run a fairly complete set of tests on your code. Being able to easily run a consistent set of tests on your code and help you quickly locate inadvertent changes in your program that might not have been caught until much later. **20 Points.**
 - k. Think about using macros in your `Makefile` for some of these.
 - l. Test early and test often.
 - m. I’m sure you will want to have additional tests; you can put them in your `Makefile` as well.
-

Things to include with the assignment (in a single `tar.bzip` file):

- 1. C source code for the solutions to the posed problems (all files).
- 2. A `Makefile` to build your code.
- 3. A text file describing what tests you ran.

Please combine all of the above files into a single `tar.bzip` file prior to submission. Run the “`make clean`” before creating the `tar.bzip` file. Consider putting a target in your `Makefile` that will run the “`make clean`” and then create the single `tar.bzip` file.